

```
In [1]: #KNN CLASSIFIER
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv("DATA/gene_expression.csv")
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Gene One	Gene Two	Cancer Present
0	4.3	3.9	1
1	2.5	6.3	0
2	5.7	3.9	1
3	6.1	6.2	0
4	7.4	3.4	1

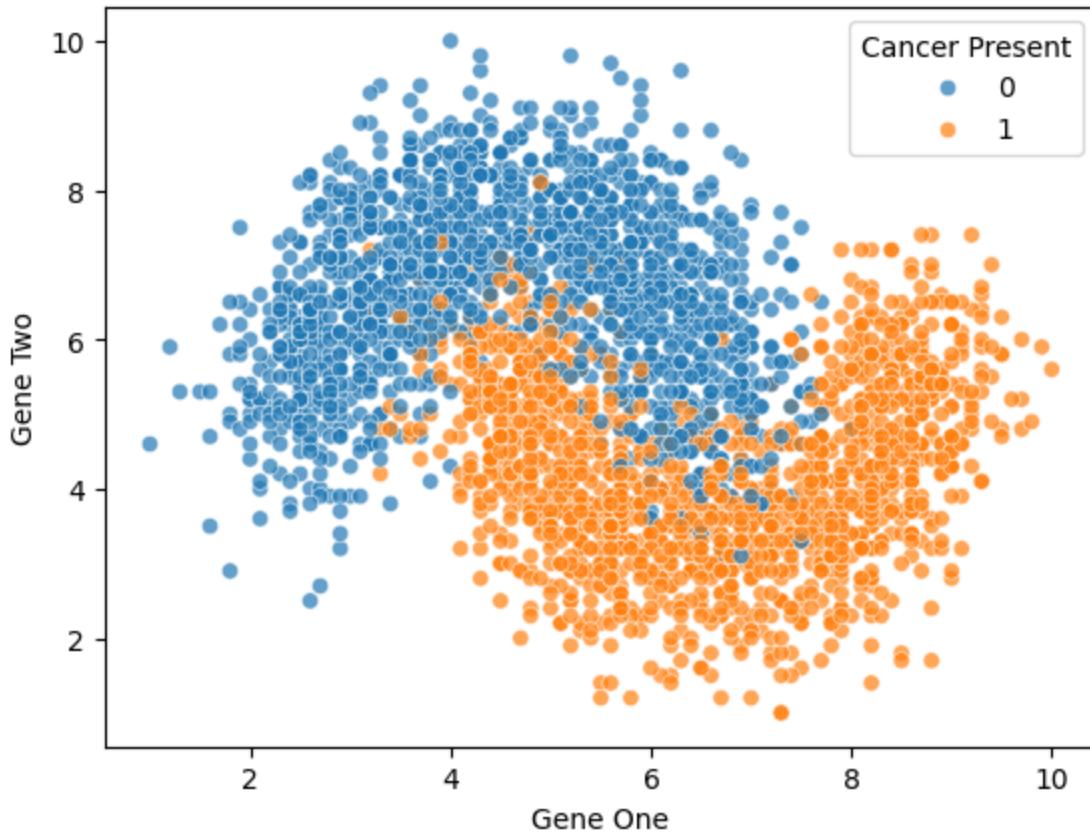
```
In [5]: df.describe()
```

```
Out[5]:
```

	Gene One	Gene Two	Cancer Present
count	3000.000000	3000.000000	3000.000000
mean	5.600133	5.410467	0.500000
std	1.828388	1.729081	0.500083
min	1.000000	1.000000	0.000000
25%	4.300000	4.000000	0.000000
50%	5.600000	5.400000	0.500000
75%	6.900000	6.700000	1.000000
max	10.000000	10.000000	1.000000

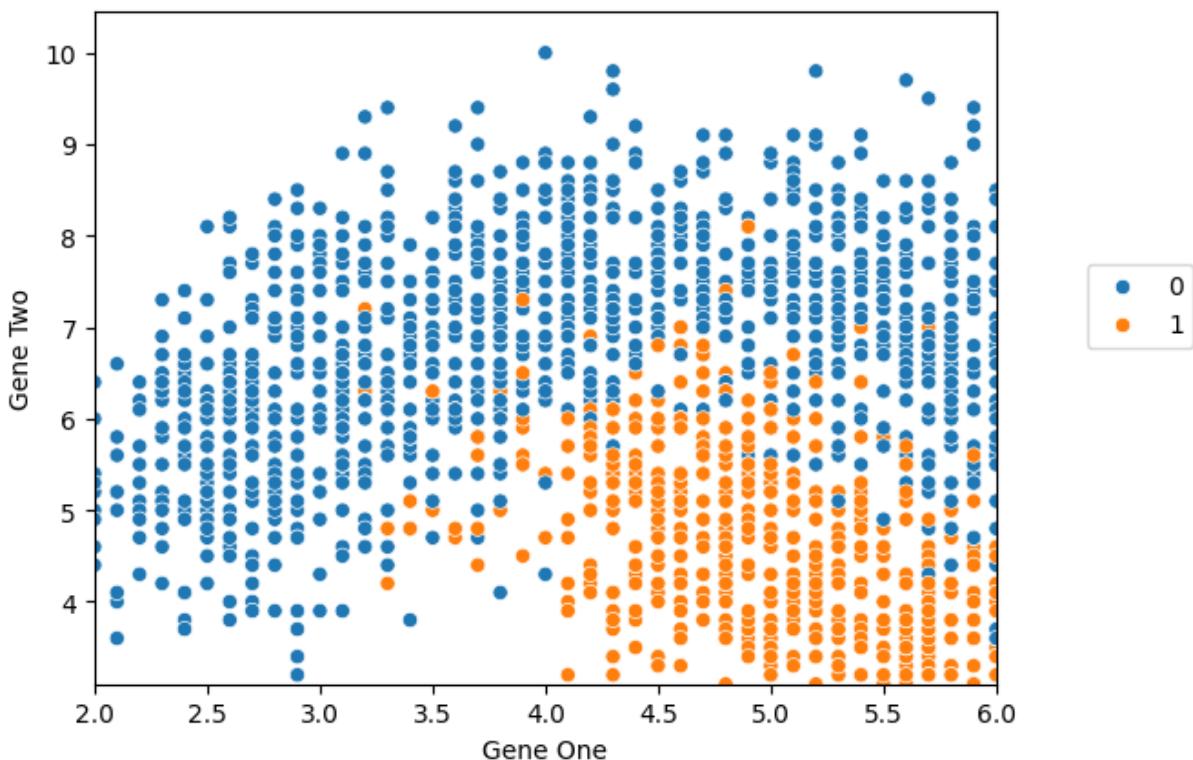
```
In [7]: sns.scatterplot(x='Gene One',y='Gene Two',hue='Cancer Present',data=df,alpha=0.7)
```

```
Out[7]: <Axes: xlabel='Gene One', ylabel='Gene Two'>
```



```
In [9]: sns.scatterplot(x='Gene One',y='Gene Two',hue='Cancer Present',data=df)
plt.xlim(2,6)
plt.ylim(3,10)
plt.legend(loc=(1.1,0.5))
```

Out[9]: <matplotlib.legend.Legend at 0x22ca303be30>



```
In [10]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
In [12]: X = df.drop('Cancer Present', axis=1)
```

```
In [13]: y=df['Cancer Present']
```

```
In [14]: X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3, random_state
```

```
In [15]: scaler=StandardScaler()
```

```
In [16]: scaled_X_train=scaler.fit_transform(X_train)
```

```
In [17]: scaled_X_test=scaler.fit_transform(X_test)
```

```
In [18]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [53]: knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(scaled_X_train, y_train)
```

```
Out[53]: ▾ KNeighborsClassifier ⓘ ?
```

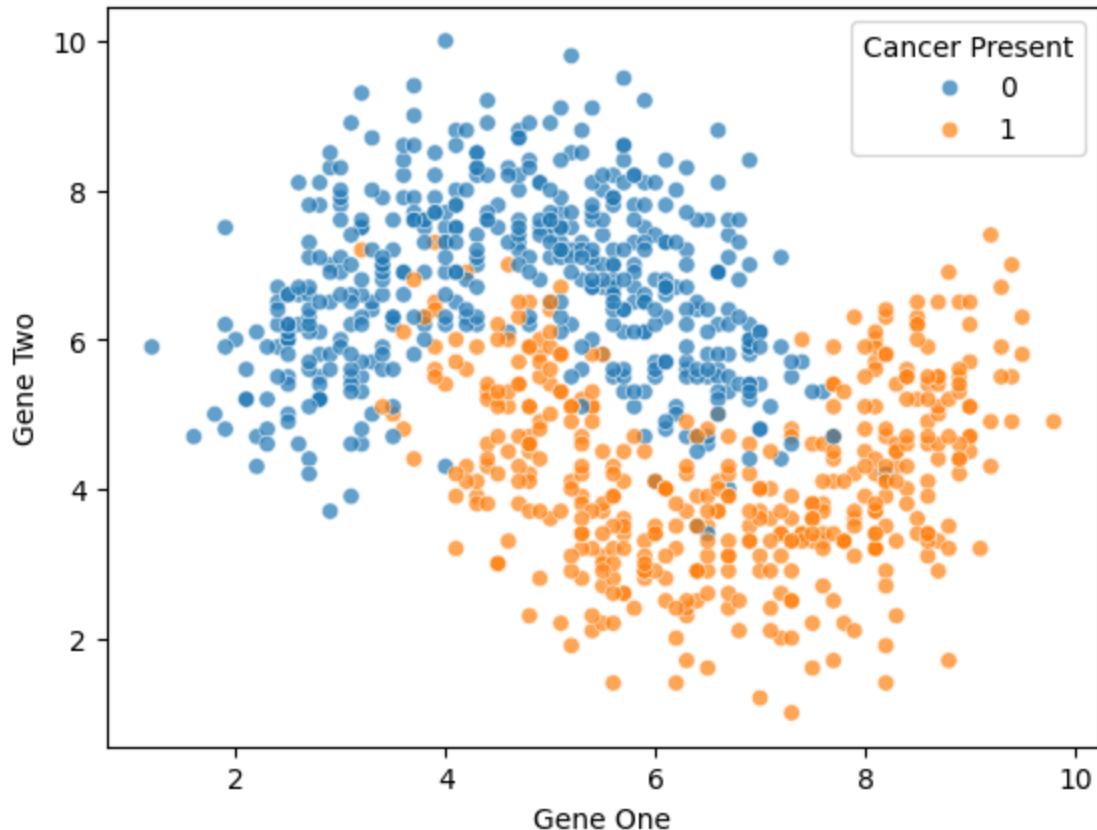
```
KNeighborsClassifier()
```

```
In [54]: full_test= pd.concat([X_test,y_test],axis=1)
len(full_test)
```

```
Out[54]: 900
```

```
In [55]: sns.scatterplot(x='Gene One',y='Gene Two',hue='Cancer Present',data=full_test,alpha
```

```
Out[55]: <Axes: xlabel='Gene One', ylabel='Gene Two'>
```



```
In [56]: y_pred=knn_model.predict(scaled_X_test)
```

```
In [57]: y_pred
```

```
In [58]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
```

```
In [59]: accuracy score(y test,y pred)
```

```
Out[59]: 0.9344444444444444
```

```
In [60]: confusion_matrix(y)
```

```
Out[60]: array([[437    33]
```

[26, 404])

... [REDACTED] ... [View Classifications](#) | [Edit](#) | [Copy](#) | [Delete](#) | [Print](#)

	precision	recall	f1-score	support
0	0.94	0.93	0.94	470
1	0.92	0.94	0.93	430
accuracy			0.93	900
macro avg	0.93	0.93	0.93	900
weighted avg	0.93	0.93	0.93	900

```
In [31]: test_error_rates=[]
for k in range(1,30):
    knn_model = KNeighborsClassifier(n_neighbors=k)
    knn_model.fit(scaled_X_train,y_train)
    y_pred_test = knn_model.predict(scaled_X_test)

    test_error=1-accuracy_score(y_test,y_pred_test)
    test_error_rates.append(test_error)
```

```
In [32]: test_error_rates
```

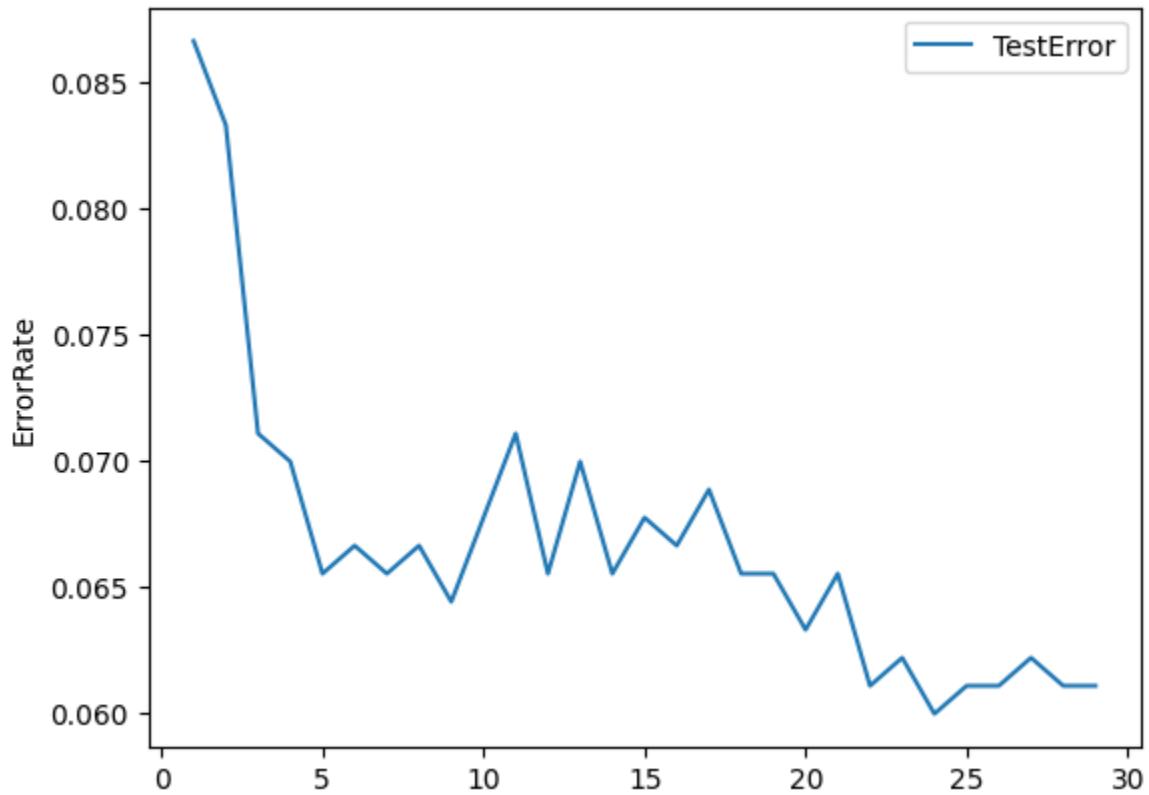
```
Out[32]: [0.08666666666666667,
 0.08333333333333337,
 0.07111111111111112,
 0.06999999999999995,
 0.06555555555555559,
 0.06666666666666665,
 0.06555555555555559,
 0.06666666666666665,
 0.06444444444444442,
 0.06777777777777783,
 0.07111111111111112,
 0.06555555555555559,
 0.06999999999999995,
 0.06555555555555559,
 0.06777777777777783,
 0.06666666666666665,
 0.06888888888888889,
 0.06555555555555559,
 0.06555555555555559,
 0.06333333333333335,
 0.06555555555555559,
 0.06111111111111116,
 0.0622222222222218,
 0.06000000000000005,
 0.06111111111111116,
 0.06111111111111116,
 0.0622222222222218,
 0.06111111111111116,
 0.06111111111111116]
```

```
In [33]: plt.figure(figsize=(10,6),dpi=200)
```

```
Out[33]: <Figure size 2000x1200 with 0 Axes>
<Figure size 2000x1200 with 0 Axes>
```

```
In [34]: plt.plot(range(1,30),test_error_rates,label='TestError')
plt.legend()
plt.ylabel('ErrorRate')
```

```
Out[34]: Text(0, 0.5, 'ErrorRate')
```



```
In [ ]:
```